

magma-vaults

November 22, 2024

Abstract

`magma-vaults` are a fork of CharmFi vaults [3], designed originally for EVM, but adapted to the Osmosis chain with minor enhancements. `magma-vaults` allow users to automatically manage up to 3 liquidity positions for Osmosis Supercharged pools [4]. A full range position, a normal concentrated and balanced one, and a third limit one using out-of-proportion balances to reduce impermanent loss. Users can configure different vaults with different parameters, make them permissionless or permissioned, and decide the exact liquidity that will go into each position.

1 Introduction

UniswapV3 introduces the concept of ticks to allow for concentrated liquidity positions over any range of ticks. Thus, we distinguish 2 spaces we want to work in: A tick space, where all ticks live, and a price space, where all prices live. As prices are logarithmic, UniswapV3 constructed a linear tick space by taking logarithms. Its in this way that UniswapV3 price function $\rho(t) = 1.0001^t$ is defined (equation (6.1) [2]), with the nice property that given the current tick t of a liquidity pool, any concentrated position equidistant to t will be balanced. That is, for any $\varepsilon > 0$, positions with tick ranges $[t - \varepsilon, t + \varepsilon]$ with reserves (x, y) will have $\rho(t) = y/x$.

One of the emergent features in UniswapV3 is that you can use concentrated liquidity positions as limit orders, with the problem that, due to their logarithmic nature, as prices grow, tick precision decreases. To give an example, if the current price of bitcoin was 80000USDC, the difference between the 2 closest ticks is of almost 8USDC. For this reason, Osmosis introduces geometric tick spacing with additive ranges [5], that allows to define an order book on top of the AMM with proper tick precision. Then, its easy to verify that Osmosis price function is:

$$\rho(t) = 10^{\lfloor t/9e6 \rfloor - 6} \left(t - 10^6 \left(9 \left\lfloor \frac{t}{9e6} \right\rfloor - 1 \right) \right)$$

Naturally, mapping the tick space and price space in this way breaks our nice equidistant property. For this reason, `magma-vaults` does operations over the price space instead of the tick space, and then takes the price function generalized inverse ρ^{-1} to decide the ticks for the balanced base position and the limit one. It is then trivial to prove that, given the current tick t and the current price $p := \rho(t)$, any position in the price range $[p/k, pk]$ will be balanced (we keep $k > 1$ for

simplicity). We refer to those factors k as `PriceFactors` in the code, and they will be, with the `Weight` type, our fundamental building block to configure custom liquidity provision strategies.

2 Architecture overview

`magma-vaults` manage up to 3 concentrated liquidity positions. A full range position, a base balanced position, and a limit position. Both, the full range position and the base position will be balanced, but as prices move, the base position could become unbalanced. Thus we expose a `rebalancing` functionality, that will burn all positions and create them again centered around the new current price. Moreover, its easy to see that the total vault reserves wont always be in proportion, so any out-of-proportion reserves will be used for the limit position. Actually, vault total reserves almost never will be in proportion, but its still theoretically possible a vault could get rebalanced to have no limit position. Thus, the core business logic of the vault contract, i.e, the rebalancing procedure, will:

1. Burn all active liquidity positions to withdraw total reserves (X, Y) . We consider new deposits to the vault to be included in this pair.
2. Calculate all the balanced reserves (x, y) in the vault. Its trivial to verify that:

$$(x, y) = \begin{cases} (X, Xp) & , Y/p > X \\ (Y/p, Y) & , \text{otherwise} \end{cases}$$

3. Calculate the reserves (x_0, y_0) to put into the full range position. For this, the vault admin (user) will decide a liquidity `Weight` $w \in (0, 1)$. For simplicity, we dont allow for extreme value of w , but the user still could make the difference between w and extremes negligible. The reserves will be calculated in such way that, if L_0 is the liquidity of the full range position, and L_1 is the liquidity of the base one, $L_0/(L_0 + L_1) = w$ will be invariant. We will elaborate on this computation in following sections.
4. Calculate trivially the reserves (x_1, y_1) to put into the base range position as $(x_1, y_1) = (x - x_0, y - y_0)$. Then, calculate the concentrated tick range as $[\rho^{-1}(p/k_1), \rho^{-1}(pk_1)]$, for current price p and `PriceFactor` k_1 decided by the user. The inverse ρ^{-1} will be discussed in following sections.

5. Calculate trivially the reserves (x_2, y_2) to put into the limit position as $(x_2, y_2) = (X - x, Y - y)$. Naturally, $x_2 = 0 \vee y_2 = 0$ is invariant, and trivial from (2). We then simply compute the tick range for the limit position as:

$$[t_a, t_b] = \begin{cases} [\rho^{-1}(p/k_2), \rho^{-1}(p)] & , x_2 = 0 \\ [\rho^{-1}(p), \rho^{-1}(pk_2)] & , y_2 = 0 \end{cases}$$

Where p is the current price, and k_2 a `PriceFactor` decided by the user.

Thus, the only remaining question is when to rebalance. For this, any vault can use up to 3 rebalancing strategies:

- `admin`: Only the vault admin can rebalance the vault whenever they feel like. Thus, its up to the admin to decide the rebalancing strategy off-chain.
- `delegate`: Any delegate address, decided by the admin, can rebalance the vault whenever they feel like. This allows for the system extensionality with, for example, oracle integration to rebalance automatically.
- `anyone`: Anyone can rebalance the vault as long as the price has moved outside the range $(p/k, pk)$ (where p is the last snapshoted price during rebalance, and k a `PriceFactor` decided by the admin) and as long as enough time has passed (threshold also decided by the admin). This allows for fully permissionless vaults if the vault admin decides to burn its ownership.

Finally, for vault withdrawals and deposits, the business logic is trivial. The only complicated part is minting vault shares (represented as a `cw20`), but for that we simply use CharmFi battle tested code.

3 Full range reserves computation

CharmFi Alpha Vaults allocate liquidity to the full range position with the formula $L = w\sqrt{XY}$ [1]. But this introduces an imprecision: As the base range size is variable, its liquidity will also be. Thus, `magma-vaults`, to make it easier to reason about fees, ensures the liquidity of the full range position remains constant through the invariant:

$$w = \frac{L_0}{L_0 + L_1}$$

Thus, for example, if we set $w = 0.5$ we will be sure that, during any period of time, both positions will earn the same amount of fees (assuming of course that the base position stays in range during that time).

Clearly, $L_0 = x_0\sqrt{p}$. We then express L_1 in terms of UniswapV3 trading curve (equation (2.1) of the paper [2]) by taking liquidity on the right, and k as the base range `PriceFactor`:

$$L_1 = x_1 \frac{\sqrt{p}\sqrt{pk}}{\sqrt{pk} - \sqrt{p}} = x_1 \frac{\sqrt{pk}}{\sqrt{k} - 1}$$

Substituting then into our w invariant:

$$\begin{aligned} w &= \frac{L_0}{L_0 + L_1} \\ &= \frac{x_0\sqrt{p}}{x_0\sqrt{p} + \frac{x_1\sqrt{pk}}{\sqrt{k}-1}} \\ &= x_0 \left(x_0 + \frac{x_1\sqrt{k}}{\sqrt{k}-1} \right)^{-1} \\ &= x_0 \left(\frac{x_0(\sqrt{k}-1) + x_1\sqrt{k}}{\sqrt{k}-1} \right)^{-1} \\ &= \frac{x_0(\sqrt{k}-1)}{x_0(\sqrt{k}-1) + (x - x_0)\sqrt{k}} \\ &= \frac{x_0\sqrt{k} - x_0}{x_0\sqrt{k} - x_0 + x\sqrt{k} - x_0\sqrt{k}} \\ &= \frac{x_0\sqrt{k} - x_0}{x\sqrt{k} - x_0} = w \xrightarrow{(A)} x_0 = \frac{\sqrt{k}wx}{\sqrt{k} + w - 1} \end{aligned}$$

Where at step (A) we simply solve for x_0 . And once we know x_0 , we of course can simply get y_0 as $y_0 = x_0p$, which ensures the reserves to be in proportion. We shall now prove that $x_0 \leq x \leq X$, as the fact will be helpful later:

$$\begin{aligned} \frac{\sqrt{k}wx}{\sqrt{k} + w - 1} \leq x &\iff \sqrt{k}w \leq \sqrt{k} + w - 1 \iff \\ \sqrt{k}(w - 1) + (1 - w) &\leq 0 \iff (w - 1)(\sqrt{k} - 1) \leq 0 \end{aligned}$$

But $(w - 1) \geq 0$ and $(\sqrt{k} - 1) \leq 0$ as $w \in [0, 1]$ and $k \in [1, \infty)$. This concludes the proof.

3.1 Computation security proof

We now prove that the way x_0 is computed in the magma-vaults codebase is secure. Thus considered the implementation:

```
pub fn calc_x0(k: &PriceFactor, w: &Weight, x: Decimal) -> Decimal {
    if w.is_zero() { return Decimal::zero() }
    do_me! {
        let sqrt_k = k.0.sqrt();

        let numerator = w.mul_dec(&sqrt_k);

        let numerator = Decimal256::from(numerator)
            .checked_mul(x.into())?; // (1)

        let denominator = sqrt_k
            .checked_sub(Decimal::one())? // (2)
            .checked_add(w.0)?; // (3)

        let x0 = numerator.checked_div(denominator.into())?; // (4)
        Decimal::try_from(x0)? // (5)
    }.unwrap()
}
```

The `do_me!` macro only creates a `anyhow::Error` closure and runs it, so we only need to ensure the commented lines won't produce any errors:

- (1). Won't overflow as $(2^{128} - 1)^2 < 2^{256} - 1$.
- (2). Won't underflow as $k \in [1, \infty)$, thus $\sqrt{k} - 1 = 0$ in the worst case.
- (3). Won't overflow as $w \in [0, 1]$, and we just subtracted 1 from \sqrt{k} .
- (4). Could only produce a division by zero if $\sqrt{k} - 1 + w = 0$. Assume $k = 1$, then $\sqrt{k} - 1 + w = 0 \iff w = 0$, in which case x_0 computation is trivial (see first line of the implementation). Note also that $k = 1 \wedge w = 0$ produces a vault with generally idle capital and only limit positions (which the first version does not support). On the other hand, the division won't overflow because $x_0 \leq x$, as stated earlier.
- (5). Finally, the downgrade to 128 bits won't fail because, again, $x_0 \leq x$, but x fits in 128 bits by definition.

4 Price function inverse

The only missing piece is the generalized inverse ρ^{-1} that will let us map every price to its closest tick in Osmosis. We will prove that such function is:

$$\rho^{-1}(p) = 10^{6-\lfloor \log p \rfloor} \left(p + 10^{\lfloor \log p \rfloor} (9\lfloor \log p \rfloor - 1) \right)$$

Of course, it could be the case that there was 2 different prices with the same tick, as we require ticks to be integers (ie, we will take $\rho^{-1}(p) := \lfloor \rho^{-1}(p) \rfloor$). Moreover, there could be 2 ticks that map to the same price, due to computational precision. That's why by generalized inverse we mean 2 things:

1. Generalized inverse property: $\forall t \in T : \rho^{-1}(\rho(t)) = t$, where $T \subset \mathbb{Z}$ is the tick space.
2. Closeness property: $\forall t_0, t_1 \in T$ s.t. $t_1 > t_0$, we have that $\forall p \in [\rho(t_0), \rho(t_1)]$ its the case that $\rho^{-1}(p) \in [t_0, t_1]$. That is, given any 2 ticks, if we map them to the price space and pick any price in that range, that price tick will also be in the tick range after pulling it back.

But its easy to check that the weaker form of the closeness property when $t_1 = t_0 + 1$ is equivalent to the general form. The generalized inverse property is not enough for our usecase, as we just mentioned there will be prices that map to the same tick. Informally, by proving the weak closeness property, we will know that any price will get mapped to its closest tick or its second closes tick (which, for our usecase, is good enough). The proof sketch goes as follows:

1. Prove the generalized inverse property by simple brute force (as the domain is finite and relatively small).
2. Observe both ρ^{-1} and ρ are monotone non-decreasing.
3. Concluded trivially that the weak closeness property holds, by simply taking inverses over $\rho(t_0) \leq p \leq \rho(t_0 + 1)$.

4.1 Computation security proof

Once again, we will have to prove that ρ^{-1} computation is secure. In the codebase:

```
let compute_price_inverse = |p: &Decimal| -> Option<i32> {  
  let floor_log_p = PositiveDecimal::new(p)?.floorlog10(); // (1)  
  let x = floor_log_p.checked_mul(9)?.checked_sub(1)?;      // (2)
```

```

let x = maybe_neg_pow(floor_log_p)? // (3)
    .checked_mul(SignedDecimal256::new(x.into()))? .ok()? // (4)
    .checked_add(p.clone().try_into().ok()?).ok()?; // (5)

let x = maybe_neg_pow(6i32.checked_sub(floor_log_p)?)? // (6)
    .checked_mul(x).ok()?; // (7)

let x: Int128 = x.to_int_floor().try_into().ok()?; // (8)
x.i128().try_into().ok() // (9)
};

```

Thus, we want to prove the closure will never return None at each step:

- (1). floorlog10 computes $\lfloor \log p \rfloor$. The implementation is:

```

pub fn floorlog10(self: &PositiveDecimal) -> i32 {
    let x: u128 = self.0.atomics().into();
    // Invariant: `u128::ilog10(u128::MAX)` fits in `i32`.
    let x: i32 = x.ilog10().try_into().unwrap();
    // Invariant: `ilog10(1) - 18 = 0 - 18` fits in `i32`.
    x.checked_sub(18).unwrap()
}

```

And as we can see, its security is already proved in the comments. Note that we subtract 18 to our result, as we compute the decimal logarithm of the atoms of a Decimal type (wrapped trivially in PositiveDecimal), thus we have to compensate for the 18 decimal points it uses. Its also trivial to verify that

$$\text{Im}(\text{floorlog10}) = [-18, 20] \subset \mathbb{Z}$$

is invariant.

- (2). Taking the upper and lower bound of the image of floorlg10, we get the valid i32:

$$(9\lfloor \log p \rfloor - 1) \in [-162, 179]$$

- (3). maybe_neg_pow simply takes SignedDecimal256::pow(10, exp) but with i32 domain instead of u32 (the implementation is trivial). Clearly, the operation could fail for the

whole domain, but in our case:

$$\begin{aligned} \text{maybe_neg_pow}(\text{Im}(\text{floorlog10})) = \\ \text{maybe_neg_pow}([-18, 20]) = [10^{-18}, 10^{20}] \subset \text{SignedDecimal256} \end{aligned}$$

Note that we're using the Signed version of `Decimal256` to compute $10^{\lfloor \log p \rfloor}$ as we're about to multiply its result by the most inner term of ρ^{-1} (that is, $9\lfloor \log p \rfloor - 1$, which we already know could be negative).

- (4). We got $(9\lfloor \log p \rfloor - 1) \in [-162, 179]$ and $10^{\lfloor \log p \rfloor} \in [10^{-18}, 10^{20}]$. Thus, the product could only overflow if $179 \cdot 10^{20}$ could not fit in `SignedDecimal256`, but that's clearly not the case.
- (5). We compute $p + x$, where x is the term computed in step (4), that is, $179 \cdot 10^{20}$ in the worst case. Thus, this addition could only overflow if `u128::MAX` + $179 \cdot 10^{20} = 2^{128} - 1 + 179 \cdot 10^{20}$ did not fit in `SignedDecimal256`, but that's clearly not the case.
- (6). We compute the left factor $10^{6-\lfloor \log p \rfloor}$. We have:

$$(6 - \lfloor \log p \rfloor) \in [6 - 20, 6 - (-18)] = [-14, 24]$$

And clearly -14 fits in `i32`, and so does 24 , and so do 10^{24} fits in `SignedDecimal256`.

- (7). Finally, we compute the product of the terms computed in (6) and (5). We know $10^{24} \cdot (2^{128} - 1 + 179 \cdot 10^{20})$ is an upper bound for our product, which already fits in `SignedDecimal256`. But note that if $\lfloor \log p \rfloor = 20$ then the left factor will be $10^{6-20} = 10^{-14}$. On the other hand, if $\lfloor \log p \rfloor = -18$, then it will be $10^{6-(-18)} = 10^{24}$. It is easy to verify then that:

$$\text{Im}(\rho^{-1}) = \left[\rho^{-1} \left(\frac{1}{10^{18}} \right), \rho^{-1} \left(\frac{2^{128} - 1}{10^{18}} \right) \right]$$

- (8). Because of (7), the integer part of p^{-1} will always fit in `Int128`.
- (9). Because of (7), the integer part of p^{-1} will always fit in `i32`.

References

- [1] Charmfi liquidity computation. <https://github.com/charmfinance/alpha-vaults-v2-contracts/blob/main/contracts/AlphaProVault.sol#L420>, 2021. Accessed: 2024-11-20.
- [2] Hayden Adams et al. Uniswap v3 core. <https://app.uniswap.org/whitepaper-v3.pdf>, March 2021. Accessed: 2024-11-20.
- [3] CharmFi. Alpha Vaults whitepaper. <https://learn.charm.fi/charm/products-overview/alpha-vaults/whitepaper>, 2021. Accessed: 2024-11-20.
- [4] Osmosis. Concentrated Liquidity module documentation. <https://docs.osmosis.zone/osmosis-core/modules/concentrated-liquidity>. Accessed: 2024-11-20.
- [5] Osmosis. Geometric Tick Spacing with Additive Ranges documentation. <https://docs.osmosis.zone/osmosis-core/modules/concentrated-liquidity/#geometric-tick-spacing-with-additive-ranges>. Accessed: 2024-11-20.